

ABHINANDAN KUMAR

Data Engineer • Snowflake • dbt • SQL • Databricks • Python

+91 98177 50262 | sabhinandan403@gmail.com | linkedin.com/in/abhinandankumar | github.com/abhinandankumar

PROFILE

Data Engineer with 3 years of experience across SQL, ETL/ELT pipelines, cloud data platforms and data-driven applications. Hands-on with Snowflake and dbt for staged-to-mart transformations, incremental processing and analytical data modeling, and with PySpark/Databricks for IoT telemetry pipelines. Strong SQL foundation across Snowflake, SQL Server and PostgreSQL, with experience supporting production data workflows, performance investigation, reporting and application integration.

TECHNICAL SKILLS

Data Platforms & Databases: Snowflake, MS SQL Server, PostgreSQL, Databricks, Delta Lake, MongoDB

Snowflake: Snowflake SQL, Stages, COPY INTO, Streams, Tasks, Query Profiling, Virtual Warehouse Sizing, MERGE Patterns, Stored Procedures

SQL: T-SQL, CTEs, Window Functions, Stored Procedures, Triggers, Views, Functions, Query Performance Tuning, Data Archival & Retention

Data Engineering: dbt, ETL/ELT, PySpark, Incremental Processing, Dimensional Modeling, Fact & Dimension Design, Data Quality

Programming & Backend: Python, SQL, Node.js, FastAPI, GraphQL, Hasura, Kafka, RabbitMQ, REST APIs, RBAC/ABAC

Cloud & Orchestration: AWS (S3, Lambda), Azure (Data Factory, PIM/RBAC), Databricks Workflows, Unity Catalog

Analytics & Tooling: Power BI, DAX, Git, Azure DevOps, Docker, TOAD, SQL Developer

AI & Emerging: LLM/agent workflows, MCP servers over data platforms, prompt engineering, OpenAI API function calling

PROFESSIONAL EXPERIENCE

Vantiva India | Software Engineer — Data Engineering

Jun 2024 – Present

Smart Spaces *Snowflake • dbt • SQL Server • PySpark • Databricks • AWS • Power BI*

- Engineered Snowflake ELT ingestion using stages, COPY INTO, Streams and Tasks; enabled scheduled incremental loads from raw/staging tables to curated analytical datasets.
- Developed dbt SQL models across staging, intermediate and mart layers, using incremental materializations and surrogate-key joins to implement reusable transformations for downstream analytics.
- Designed analytical data models by establishing source grain, business keys and fact-to-dimension relationships; built curated marts for KPI reporting and operational analysis.
- Diagnosed failed Snowflake loads and scheduled Tasks through load, task and query history; investigated schema and file-format errors, validated reruns, and used query profiles and warehouse metrics to tune recurring workloads.
- Authored SQL Server stored procedures, triggers and views; developed complex T-SQL for multi-table joins, CTE-based transformations and reporting workflows, while supporting historical-data archival and retention.
- Built PySpark DataFrame-based Databricks ETL pipelines to cleanse, transform and aggregate 10,000+ daily IoT telemetry points, automating Wi-Fi mesh uptime calculations.
- Implemented a serverless AWS Lambda reporting workflow that generated parameterized Excel reports across 50+ sites; used SQL analysis to trace LoRa/Zigbee connectivity and missing-packet issues, helping maintain a ~99% uptime SLA.
- Built Power BI dashboards used by sales teams, senior leadership and clients to monitor device performance and usage statistics.

HomeSight Care *Node.js • Kafka • In-memory Caching • RBAC/ABAC*

- Migrated core APIs from direct database reads to a Kafka-backed in-memory caching layer with indexed lookups and database fallback, reducing high-traffic response times from ~7 seconds to 40–50 ms.
- Integrated RBAC/ABAC controls across user-module APIs to enforce consistent authorization for personal-data and hub-information workflows.

Catalyst Data Lake — Platform Access Engineering *Azure • ARM RBAC • Azure PIM*

- Implemented just-in-time (JIT) privileged access for production ARM RBAC roles on the Catalyst data lake platform, working through documented remediation steps with a review-before-execution approach.

Wow Labz | Backend Engineer Intern

Apr 2024 – Jun 2024

Cab Booking Application & Video Processing Pipeline *Python • FastAPI • MongoDB • Kafka • AWS S3*

- Built a FastAPI reporting service used daily by administrators and 10–12 drivers to track driver performance, revenue and bookings.
- Used Kafka-based messaging to move video data through a seven-stage pipeline for chunking, transcription, subtitles and multilingual dubbed audio.

Yamaha Motor Solutions India | Graduate Engineer Trainee — Full Stack Developer

Jul 2023 – Apr 2024

Peer Review Application *PostgreSQL • GraphQL • Node.js • RabbitMQ*

- Built a peer-review application adopted by a 30-engineer team, replacing manual review tracking with automated status visibility and real-time RabbitMQ notifications.
- Developed GraphQL APIs with Hasura on PostgreSQL; authored complex SQL functions, cursors and dynamic DML queries for application workflows.

CERTIFICATIONS

- Databricks Certified Data Engineer Professional; Databricks Fundamentals Accreditation
- McKinsey Forward Learning Program — Problem Solving, Leadership & Communication

EDUCATION

B.Tech, Computer Science Engineering — Kurukshetra University | GPA 8.02 / 10 2019 – 2023